Claude 4 Opus chat on:

User:

metadata: BB

Hi Claude. I would like you to help us construct a prompt (system and however many necessary user prompts) for the following goal.

We are building a boolean index on a very large corpus. Specifically, we want each category of that index to indicate if the page speaks to a given topic or not. We need to turn these general research questions into very specific deductive research questions that can be answered by:

1. **Dropping the appropriate system and preparatory prompts through a LLM**
2. **With that setup and ollama, take the deductive questions and the the LLM to output blockquotes if there is a matching phrase which responds to the question with the right semantic feel. (We will note this archive covers two decades of content, and we will want to make sure the AI is in the right evaluative accent to interpret the page in its own context.**
3. **If there is a blockquote output AND that blockquote can be found in the source text by python, then, emit the blockquote to a file.**

For reference, here is what my students did for an ancient history topic.

The research questions we would like to operationalise for this boolean index are:

> **How are children addressed on these websites (by addressed, I mainly mean choice of words, what actions they are invited to take, what topics they are encouraged to discuss, etc.)? How does that change over time?**

> **How are girls and boys addressed in the corpus?**

> **What are common features in terms of design across the websites? How does that develop over time?**

> **How did kids contribute to these websites?**

We need to break these into specific categories that can be pulled out, one at a time, by the LLM.

To build these prompts, it is essential that you ask us one critical, specific, and pointed question at a time until you feel that you understand our disciplinary context, our goals, and what we want out of this prompt structure sufficiently to build our prompts with us.

It is important that you are critical and specific, not complimentary and positive. We are academics, we already know our ideas are good.

To begin, functionally decompose the task at hand, and give me a readback about the communication styles we require.

Assistant:

Functional Decomposition

Your task requires transforming broad historiographical questions into discrete, binary classification problems. The pipeline involves:

1. **Prompt Engineering**: Creating system prompts that establish disciplinary context and evaluative stance appropriate to temporal periods within your corpus
2. **Semantic Query Design**: Converting research questions into specific, testable propositions that yield extractable evidence
3. **Evidence Extraction**: LLM identifies and outputs verbatim blockquotes matching semantic criteria
4. **Verification**: Python string-matching confirms blockquote existence in source
5. **Boolean Classification**: Presence of verified blockquote = True for that category

Communication Requirements

You require:

- Direct, critical questioning without performative courtesy
- Academic Australian English (honour, analyse, etc.)
- Specificity over generality
- Decomposition of complex phenomena into observable textual features
- Temporal sensitivity in interpretation

Critical Question 1

Your corpus spans "two decades" of web content about/for children. Which two decades specifically? The interpretive framework for evaluating how children are "addressed" in 1995 HTML differs substantially from 2015 responsive design. The semantic markers of "child-directed speech" in early web forums versus social media platforms require different evaluative criteria.

What is the temporal range of your corpus, and what types of web content does it contain (static pages, forums, blogs, social platforms)?

User:

metadata: BB

Assistant:

Corpus Clarification

1996-2005, ages 5-15. Forums largely absent due to password protection.

Critical Question 2

You mention "books for kids" and "form-ish" responses. What precisely constitutes your corpus? Are these:

- Corporate websites targeting children (Disney, Nickelodeon, educational publishers)?
- Individual/amateur sites created for children (teachers' pages, parent resources)?
- Sites created BY children (personal homepages, school projects)?
- Mixed authorship (adult-moderated spaces with child contributions)?

The semantic markers for "how children are addressed" differ fundamentally between corporate pedagogical content versus peer-to-peer communication. A Disney.com game instruction addresses children differently than a 12-year-old's Geocities fan page addresses visitors.

Additionally, what proportion of your corpus contains actual child-generated content versus adult-generated content about/for children? Your fourth research question implies child contributions exist, but the absence of forums suggests limited interactive content.

User:

metadata: BB

****Brian:** [00:00] You mentioned books for kids and form-ish responses. I think we meant forum-ish responses. What precisely constitutes your corpus? So, for sake of argument, our full corpus is going to be much larger than our test corpus. Um so do No, let's let's answer for our full corpus. **Helle:** [00:24] Our like full corpus in the research? **Brian:** [00:26] No no, as in like the test corpus. **Helle:** [00:28] The the non-toy test corpus that we'd be running this over if this worked to scale. **Brian:** [00:34] All right. **Helle:** [00:35] Then it spans a lot of different kinds of websites. There are like corporate websites targeting children. Uh there are individual amateur sites created for children, uh by teachers, uh not by parents but also uh amateur websites um by like fan web pages, so like with Buffy the Vampire Slayer um or Star Wars. Harry Potter. um There are sites created by children or uh in an infrastructure like either by children or in an infrastructure created by adults, but then populated by children. **Brian:** [01:15] Yep. **Helle:** [01:16] Um and uh yeah, so also personal home pages by children. Um there yes, and mixed authorship. **Brian:** [01:24] Yeah. And we are actually quite interested in how how these different authors swap register to speak to different, to speak to children because the regist as you note, register of Disney is going to be very different than a register of a 12-year-old Geocities. And so we are very much interested in I don't I'd probably actually say author register as one of these like is this a corporate web page? Was this web page, was this source written by sorry, no. Does this page have source contributed from a corporate, from a teacher, from a child, et cetera? So we actually care about each one of those bullets as one of the categories that we would like our Boolean index to pull out. Yes. And we have no idea about proportion, that's why we're doing this research area. **Helle:** [02:26] Yep. **Brian:** [02:27] So, who knows? That's what research is for.**

Assistant:

Authorship Categories Identified

Corporate, teacher-created, fan sites, child-created, mixed adult-child. Register variation is central.

Critical Question 3

How will you operationalise "child-addressed" content versus incidental child presence? A Buffy fan page might contain:

- Direct address to young fans ("Hey kids, check out these cool Buffy pics!")
- Peer-to-peer child communication ("im 11 and i think spike is cute")
- Adult discussion about child-appropriate content ("This episode might be too scary for younger viewers")
- Child-generated content not addressing children ("My essay on symbolism in Buffy")

Your first research question asks how children are "addressed" - but distinguishing deliberate address from mere topical relevance requires specific linguistic markers. Do you have ground-truth examples of what constitutes "addressing children" in your corpus?

For instance, is imperative mood directed at assumed child readers ("Click here!" "Join our club!") sufficient? Or do you require explicit age-markers ("For kids 8-12")? What about simplified vocabulary without explicit child-targeting?

This distinction determines whether your boolean categories should focus on linguistic features (imperative verbs, exclamations, simplified syntax) or semantic content (references to school, parents, age-appropriate activities).

User:

metadata: BB

Assistant:

Classification Schema Clarified

Trinary (yes/maybe/no) for child-directed content. No ground truth. Seeking discovery through broad categorisation.

Critical Question 4

Your blockquote extraction mechanism requires exact string matching. How will you handle:

1. **OCR errors** from scanned content (if any)?
2. **HTML markup** within text spans? ("Click here for fun!" vs "Click here for fun!")
3. **Non-standard characters** from 1990s encoding issues (curly quotes, em-dashes, character replacements)?
4. **Whitespace variations** (multiple spaces, tabs, linebreaks)?

The LLM will output clean text, but your source files might contain ` `, `<br>`, or Windows-1252 encoded characters. Will you pre-process sources to plaintext, or must the LLM output preserve HTML entities?

Additionally, what constitutes a valid "blockquote" length? Single imperative sentences ("Join us!")? Full paragraphs? The granularity affects both extraction reliability and analytical utility. A boolean index saying "contains child-directed imperatives" differs from one preserving "Join our Pokemon club today, trainers aged 8-12!"

User:

metadata: BB

****Brian:** [00:01] So, will the LLM So, the LLM will output clean text, but our source files might be awful. **Helle:** [00:07] Yeah. **Brian:** [00:10] I think I'm going to be dropping these pages through Pandoc today. Actually, no, I will be experimenting with whether or not cleaning them down to Markdown improves the quality of responses. So, the answer is, we may I'll be running it twice. Four times, actually, cause I'll be running on two different models twice. So, we might preprocess via Pandoc. We, well, that is one of our variables, and then the output does not need to preserve HTML. Um a valid block quote is Yeah, single and imperative sentences are just fine. Um and no, child-directed imperatives is much more correct than that. And actually, quite frankly, because we can be fancy, putting the most representative CSV, er its uh um block quote in the CSV, rather than yes, maybe, no, is much better because then we can just scan it and see what it's there. So, therefore, we care about the CSV from a presence absence perspective, and it will put block quotes in if there's anything that's matching. Um and fuzzy match. So, so the block quote matching should be fuzzy match to the source text, assuming that there's going to be HTML in the way. Um which means we need to block quote match using Beautiful Soup. **Helle:** [01:52] Okay. **Brian:** [01:54] Yay**

Assistant:

Technical Pipeline Clarified

Pandoc preprocessing (experimental), fuzzy matching via Beautiful Soup, representative blockquotes in CSV rather than boolean values alone.

Critical Question 5

Your second research question targets gendered address ("How are girls and boys addressed?"). This requires distinguishing:

1. **Explicit gendering** ("Hey boys!" "Girls will love this!")
2. **Pronoun assumptions** ("Ask your mum" vs "Dad can help")
3. **Activity-based gendering** (dolls/football without explicit gender markers)
4. **Visual gendering** (pink backgrounds - but you're processing text)

Given HTML/Markdown input without image analysis, will you rely solely on textual gender markers? Or are you inferring gender targeting from semantic domains (princess games = girls)?

More critically: How do you distinguish pages addressing "boys" from pages merely mentioning boys? A page about "Books for Boys" addresses parents/teachers, not boys themselves. Do you want both, or only direct address TO gendered children?

Finally: Non-binary/gender-neutral address ("kids," "friends," "trainers") - separate category or ignored?

User:

metadata: BB

****Brian:** [00:00] Uh one thing to note is that this is going to be multilingual. This is going to be English, Danish, and Korean. Um and yes, we want to differentiate explicit gendering, assumptions, activity gendering, and visual gendering. Um so, each one of those is a category. Yes, each one of those is absolutely a category. Um and and we're going to see whether or not the LLM can parse HTML into some sort of thing where it can make gender-based statements. **Helle:** [00:43] Oh boy. If possible. **Brian:** [00:46] If possible. Um how do you differentiate or distinguish pages addressing boys from pages mentioning boys? Um **Helle:** [01:01] well, it's crucial that it doesn't like distinguish things that mentioning boys as for boys, because girls could be talking about boys. **Brian:** [01:09] Yes. So, do we care about a page giving information about children to parents? **Helle:** [01:22] Um well, I think the most important thing here is actually something different, which is I think most pages were um for kids in general. **Brian:** [01:32] Yeah. **Helle:** [01:33] But um in the corpus, it is expected that there are separate pages for girls only. **Brian:** [01:41] So so therefore, is this page intended for a specific gender? **Helle:** [01:48] Yeah. **Brian:** [01:49] Is this page intended for kids, implicitly boys, but **Helle:** [01:53] yeah. **Brian:** [01:55] yeah, yeah. And then and I would say a page about books for boys, the judgment call is register. **Helle:** [02:01] Yeah. **Brian:** [02:01] Like, if if it is talking about books that your children will enjoy, that is different than talking about books you will enjoy. **Helle:** [02:13] Yeah. **Brian:** [02:14] And so we need to differentiate address so that children are involved in the communication cycle, not just the topic cycle. **Helle:** [02:25] Yes. Yeah. Yeah. That's it. **Brian:** [02:31] Um and yes, um uh uh non-binary and gender-neutral address should be its own category.**

Assistant:

Multilingual Gender Categories Confirmed

English, Danish, Korean. Four gender-address types plus gender-neutral. Children must be in communication cycle, not merely topical.

Critical Question 6

Your third research question concerns "common features in terms of design." From HTML/Markdown, you can extract:

1. **Structural markers**: Frame usage, table layouts, image-to-text ratios
2. **Interactive elements**: Form presence, JavaScript alerts, button text
3. **Typographic indicators**: Font tags, colour specifications, text sizing
4. **Navigation patterns**: Menu structures, link density, site depth

However, "design" in 1996-2005 often meant inline styling, font tags, and bgcolor attributes - not CSS. Will your LLM parse raw HTML to identify `<font color="red">`, `<blink>`, `<marquee>`? Or are you inferring design intent from structural patterns?

More precisely: Are you seeking technological affordances (presence of forms = interactivity), aesthetic choices (colour words in HTML), or information architecture (navigation complexity)? Each requires different extraction strategies.

Given your temporal interest, should categories distinguish early HTML practices (table-based layouts) from later CSS adoption? Or is "colourful design" equally valid whether achieved via `bgcolor="#FF69B4"` or inline styles?

User:

metadata: BB

Assistant:

Design Categories Established

JSON output with multiple blockquotes. Technological affordances (yes/maybe/no), aesthetic choices (colour/font listing), complexity inference. Viewer-received aesthetics over implementation details.

Critical Question 7

Your fourth research question - "How did kids contribute?" - requires distinguishing child-generated content from adult content. In 1996-2005, what constitutes reliable evidence of child authorship?

Consider:

- **Orthographic markers**: Spelling errors, but which types? Phonetic spellings ("kool"), age-appropriate mistakes ("there/their"), or keyboard slips?
- **Syntactic patterns**: Run-on sentences, fragment sentences, inconsistent capitalisation?
- **Semantic markers**: Topic mixing ("my pokemon cards and my sister is annoying"), school-grade references?
- **Performative age claims**: "im 12" - but adults posing as children existed even then

Without metadata, how will you weight self-reported age ("I'm 10") against linguistic evidence? A teacher writing simplified content for children might trigger false positives. Conversely, a precocious 12-year-old might write "properly."

Should the system distinguish probable child authorship from confirmed child spaces (explicitly child-submitted sections)? Or combine all evidence into single confidence scores?

What about collaborative content where children and adults both contributed?

User:

metadata: BB

****Brian:** [00:09] Topic mixing? **Helle:** [00:11] Yes. **Brian:** [00:19] Well, and so the answer, so without metadata, the whole point of this is presence or absence without judgment. So, are there orthographic markers? Are there syntactic patterns? Are there semantic markers? Are there performative age claims? Yes, ASL and false ASL responses has been present since UseNet. What we care about the AI doing is merely extracting those claims. Oh. Uh no, it should make no judgment calls about authorship. If there is a plausible child involved either as audience or as author, it should match.**

Assistant:

Child Involvement Markers

Extract evidence without authorship judgment. Any plausible child involvement (audience/author) triggers match.

Critical Question 8

You mentioned the LLM needs "the right evaluative accent to interpret the page in its own context" across 1996-2005. How granular should temporal contextualisation be?

Consider "cool" - universally positive across your period. But "phat" (1996), "tight" (1999), "sick" (2003) carry temporal signatures. Should your system:

1. **Ignore temporal variation**: Treat all positive youth slang as equivalent?
2. **Flag anachronisms**: Note when 2004 slang appears in supposedly 1996 content?
3. **Adjust interpretation**: Recognise "sick website" as positive post-2000 but potentially negative in 1996?

Related: Technology literacy assumptions shift rapidly. "Click here" is redundant by 2005 but necessary scaffolding in 1996. "Ask your parents before downloading" changes meaning as broadband replaces dial-up.

Should the system prompt include year-specific contextual knowledge? Or one unified prompt treating the decade as culturally homogeneous?

How do you want to handle multilingual temporal markers - Danish/Korean youth slang evolution likely differs from English patterns?

User:

metadata: BB

****Brian:** [00:02] No, it should not engage. No, um it should ignore temporal variation. It should just be primed to recognize it. **Helle:** [00:18] Yeah. It doesn't it doesn't know that. **Brian:** [00:20] Yeah. It doesn't know that. More to the point, it's really LLMs are very bad with temporal judgment anyways. So, no, we should not include year specific contextual knowledge. We should not do anything other than saying we are looking for by at utterances and performances by children in this range. That's it. **Helle:** [00:43] Yeah. This absolutely not. **Brian:** [00:45] Yeah. No judgment at all here. And we are going to ignore temporal variation. Anything that could be construed in context as youth slang in any of these languages should be valid for these. **Helle:** [00:59] Yeah.**

Assistant:

Temporal Approach Confirmed

No temporal judgments. Recognise all potential youth markers across decade uniformly. Extract evidence without evaluation.

Critical Question 9

Given JSON output with multiple blockquotes per category, how should the LLM handle conflicting evidence within a single page?

Example: A teachers' resource page might contain:

- Adult-to-adult instruction ("This worksheet helps students...")
- Adult-to-child materials ("Hey kids, let's learn about fractions!")
- Sample child responses ("i think fractions r hard but fun")

Should the system:

1. **Output all matches**: Every blockquote for every category triggered, even if contradictory?
2. **Hierarchical classification**: Primary audience overrides embedded examples?
3. **Section-based extraction**: Treat visually/structurally separated sections independently?

This particularly matters for pages with multiple registers - a forum index showing both moderator posts and child replies, or educational sites with parallel parent/child sections.

Additionally: What's your maximum blockquote length? Full paragraphs might provide context but reduce scanability. Single sentences might lose critical nuance. Should the system extract minimal sufficient evidence or comprehensive passages?

User:

metadata: BB

****Brian:** [00:02] Well, so remember, each category is going to be within its own context. Each instance of the LLM is not going to be aware of other categories. So it should just output all all viable matches. **Helle:** [00:19] I cannot I cannot see the consequences of this. I'm so sorry. Like I cannot… **Brian:** [00:24] Oh no. This this is something we're going to need to tweak. Like right now, we're going to start by outputting all viable matches, which means that we're going to cast a broader net, which is more important than casting a narrower net. **Helle:** [00:37] Yeah. **Brian:** [00:38] Uh maximum block quote length, uh paragraph. **Helle:** [00:42] Yeah.**

Assistant:

Processing Approach Clarified

Each category runs independently. Output all viable matches. Paragraph maximum. Broad net prioritised over precision.

Critical Question 10

You're building a boolean index with independent category assessments. How many distinct categories do you envision total?

Current implied categories include:

- Direct child address (yes/maybe/no)
- Corporate authorship
- Teacher authorship
- Child authorship markers (orthographic/syntactic/semantic/performative)
- Gender targeting (explicit/pronoun/activity/visual/neutral)
- Interactive affordances
- Aesthetic patterns
- Complexity assessment
- Language-specific markers (English/Danish/Korean)

That's potentially 15-20+ categories per page. Will you run 20 separate LLM instances, or batch related categories?

More critically: Where's the boundary between evidence extraction and interpretation? "Contains imperatives" is extraction. "Imperatives targeting children" requires interpretation. Your "yes/maybe/no" framework implies interpretation, but you emphasise extracting without judgment.

Should prompts request: "Find text that might indicate child address" (interpretive) or "Find imperatives, exclamations, and second-person pronouns" (mechanical)?

User:
metadata: BB

let's go for 15 well operationalised categories per page. And it's just going to be in a for loop. No/minimal interpretation. If there is a category that requires interpretation, either we extract evidence that can be extracted to support *Human* interpretation or we skip it. The more clean, specific, and deductive the AI's request, the better.

Assistant:

Prompt Architecture

Based on our discussion, here's the proposed prompt structure for your boolean index system:

System Prompt

User:

You are a text extraction tool for historical web research. You extract verbatim blockquotes from 1996-2005 web pages that contain specific linguistic or structural features.

You make no interpretive judgments about intent or audience. You only identify and extract text containing specified patterns.

When you find matching text, output it exactly as it appears in the source, preserving original spelling, capitalisation, and errors. Maximum blockquote length: one paragraph.

Output format: JSON with three possible values for each query:
- "match": "yes" with "blockquotes": [array of matching text]
- "match": "maybe" with "blockquotes": [array of possible matches]
- "match": "no" with "blockquotes": []

Context: These pages may contain content in English, Danish, or Korean for or by people aged 5-15. Extract evidence regardless of language.
```

### Category-Specific Prompts (15 operationalised categories)

**Category 1: Imperative Verbs**
```
Extract all sentences containing imperative verbs (commands/instructions).
Examples: "Click here", "Join us", "Download now", "Klik her", "클릭하세요"
Include exclamations with implied commands: "Let's play!", "Time for fun!"
```

**Category 2: Explicit Age/Child References**
```

```
Extract text containing age numbers 5-18, or words: kids, children, boys, girls, børn, barn, 아이들, 어린이
Include: "for ages", "år", "살", grade levels, "graders"
```

Category 3: Corporate Register Markers
```
Extract text containing: "our products", "company", "corporation", "trademark", "©", "all rights reserved", "privacy policy", "terms of service"
Include formal business language: "we provide", "our mission", "virksomhed", "회사"
```

Category 4: Educational Register Markers
```
Extract text containing: "students", "classroom", "homework", "assignment", "worksheet", "learn", "teacher", "elever", "lærer", "학생", "선생님", "숙제"
Include: "parents", "guardians", "forældre", "부모님"
```

Category 5: Non-standard Spelling
```
Extract text with:
- Number substitutions: "2" for "to", "4" for "for", "u" for "you"
- Phonetic spellings: "kool", "wuz", "cuz", "plz"
- Multiple repeated letters: "coooool", "yesssss"
- Mixed case in single words: "PoKeMoN", "SpOnGeBoB"
```

Category 6: Syntactic Irregularities
```
Extract sentences that:
- Lack capitalisation at start
- Lack ending punctuation
- Contain multiple punctuation marks: "!!!", "???", "!?!?"
- Are fragments without subjects or verbs
```

Category 7: Topic Mixing Markers
```
Extract text containing sudden topic shifts within single passages.
Look for connecting words used incorrectly: "and", "but", "so" joining unrelated ideas
Example patterns: "my dog is brown and tomorrow is saturday and i like pizza"
```

Category 8: Age/Identity Claims

```
Extract text containing: "I am [number]", "I'm [number]", "years old", "år gammel", "살이에요"
Include: "my age", "we are", "grade", "klasse", "학년"
Include abbreviations: "m/f/[number]", "asl"
```

Category 9: Gendered Direct Address
```
Extract text containing direct address with: "boys", "girls", "drenge", "piger", "소년", "소녀"
Include: "ladies", "gentlemen", "guys", "dudes"
Must be addressing reader, not describing: "Hey boys" not "books about boys"
```

Category 10: Gendered Activities/Objects
```
Extract text referencing:
- Dolls, princesses, makeup, ballet, pink
- Trucks, soldiers, football, weapons, blue
- "Dukker", "prinsesse", "fodbold", "인형", "공주", "축구"
Only when presented as activities/interests, not in other contexts
```

Category 11: Interactive Element Text
```
Extract text from or near interactive elements:
- Submit/button text: "submit", "send", "go", "enter"
- Form labels: "name:", "email:", "your"
- JavaScript alerts or prompts in HTML
Include Danish: "send", "navn" and Korean: "제출", "이름"
```

Category 12: Colour/Style Words in HTML Context
```
Extract HTML attributes or text describing colours:
- Color names: "red", "blue", "rainbow", "rød", "blå", "빨간"
- HTML: "bgcolor=", "color:", "<font color="
- Descriptive: "colourful", "bright", "farverig", "화려한"
```

Category 13: Youth Slang/Informality
```
Extract informal expressions:
- Intensifiers: "totally", "super", "mega", "way cool"
- Positive slang: "awesome", "rad", "sick", "tight", "phat", "fed", "sejt", "대박"
- Discourse markers: "like", "you know", "whatever"
```

- Emoticons: ":)", ":(", "^_^"
```

**Category 14: Question Forms**
```
Extract direct questions to reader:
- "Do you", "Have you", "What's your", "Can you"
- "Har du", "Kan du", "너는", "있니"
- Includes rhetorical questions and quiz questions
```

**Category 15: First-Person Plural Inclusivity**
```
Extract "we", "us", "our", "let's", "vi", "os", "우리"
When used to include reader: "We can play", "Let's explore"
Not corporate we: distinguish "We love games!" from "We are a company"
```

### Processing Instructions

For each category, run independently:

1. Scan entire document for pattern matches
2. Extract all matching passages (up to paragraph length)
3. If clear matches exist: "yes"
4. If ambiguous/partial matches: "maybe"
5. If no matches: "no"
6. Output JSON with all blockquotes found

Would you like me to refine any of these categories or adjust the extraction criteria?

---

# User:

**metadata: BB**

**Give me a comprehensive and exhaustive summary of this conversation, our decisions, and our intentions. I need to move this chat to a project and cannot simply move this specific context.**

---

# Assistant:

# Comprehensive Summary: Boolean Index for Children's Web Corpus (1996-2005)

### Project Overview

You're building a boolean index system to analyze a decade of web content (1996-2005) targeting or involving children aged 5-15. The corpus is multilingual (English, Danish, Korean) and spans diverse sources from corporate websites to child-created Geocities pages.

### Core Research Questions

1. **How are children addressed on these websites?** (word choice, calls to action, topics encouraged)
2. **How are girls and boys addressed differently?**
```

3. **What are common design features across sites?** How do they develop over time?
4. **How did kids contribute to these websites?**

Corpus Composition

- **Corporate websites** targeting children (Disney, educational publishers)
- **Individual/amateur sites** for children (teacher pages, fan sites for Buffy/Star Wars/Harry Potter)
- **Child-created content** (personal homepages, contributions to adult-created infrastructures)
- **Mixed authorship** spaces
- **Notable absence**: Password-protected forums (not archived)
- **Temporal markers**: 1997 Korean economic crisis, 9/11, emergence of Web 2.0 (~2002)

Technical Pipeline

1. **Pre-processing**: Experimental Pandoc conversion to Markdown (testing both raw HTML and cleaned versions)
2. **LLM Processing**: Each category runs independently through the LLM (15 separate runs per page)
3. **Extraction**: LLM outputs JSON with blockquotes for matching content
4. **Verification**: Python fuzzy-matching via Beautiful Soup to confirm blockquotes exist in source
5. **Output**: JSON format with yes/maybe/no classifications plus extracted blockquotes

Critical Methodological Decisions

Interpretation vs Extraction

- **NO interpretation** of intent or audience by the LLM
- Extract evidence for **human interpretation**
- Cast a **broad net** - over-inclusion preferred to missing data
- Each category assessed **independently** (no cross-category awareness)

Temporal Approach

- **NO temporal judgments** or year-specific contextualisation
- Treat entire decade as single period
- Recognise all potential youth markers uniformly across timespan
- Ignore anachronisms or slang evolution

Classification Schema

- **Trinary system**: yes/maybe/no for each category
- **All matches extracted** even if contradictory within a page
- **Paragraph maximum** for blockquote length
- Output **all viable matches** rather than most representative

Child Involvement

- Extract markers **without authorship judgment**
- Any plausible child involvement (audience OR author) triggers match
- No distinction between genuine child content and adult-created simplified content
- Include direct address, peer communication, and implicit targeting

Gender Analysis

- Track explicit gendering, pronoun assumptions, activity-based gendering, visual gendering (if parseable from HTML)
- Distinguish pages addressing gendered children from pages merely mentioning them
- Children must be in **communication cycle**, not just topical
- Include gender-neutral/non-binary address as separate category

15 Operationalised Categories

1. **Imperative Verbs** - Commands and instructions
2. **Explicit Age/Child References** - Direct mentions of ages 5-18 or child-related terms
3. **Corporate Register Markers** - Business language, copyright notices
4. **Educational Register Markers** - School-related terminology
5. **Non-standard Spelling** - Phonetic, number substitutions, repeated letters
6. **Syntactic Irregularities** - Fragments, missing punctuation/capitalisation
7. **Topic Mixing Markers** - Sudden shifts within passages
8. **Age/Identity Claims** - Self-reported ages, ASL patterns
9. **Gendered Direct Address** - "Hey boys/girls" addressing readers
10. **Gendered Activities/Objects** - Dolls/trucks when presented as interests
11. **Interactive Element Text** - Button labels, form text
12. **Colour/Style Words in HTML** - Aesthetic descriptions and HTML attributes
13. **Youth Slang/Informality** - Period-appropriate informal expressions
14. **Question Forms** - Direct questions to readers
15. **First-Person Plural Inclusivity** - "We/us/let's" including reader

Key Principles

- **No ground truth** - This is exploratory research to discover patterns
- **Register awareness** - Differentiate corporate, educational, peer, and child voices
- **Multilingual equality** - Extract evidence regardless of language
- **Design as received** - Care about viewer experience, not implementation (table-based vs CSS)
- **Communication focus** - Children must be addressed/involved, not just discussed
- **Evidence over inference** - When in doubt, extract textual evidence rather than interpret

Technical Specifications

- **Output format**: JSON with nested blockquote arrays
- **Matching**: Fuzzy matching to handle HTML artifacts, encoding issues, whitespace variations
- **Processing**: Simple for-loop through categories (no batching)
- **Models**: Testing on two different LLMs with both raw and cleaned HTML
- **Languages**: Must handle English, Danish, and Korean equivalents for all categories

This system prioritises comprehensive evidence collection over precision, allowing human researchers to interpret patterns across the corpus rather than relying on AI interpretation of historical children's web culture.

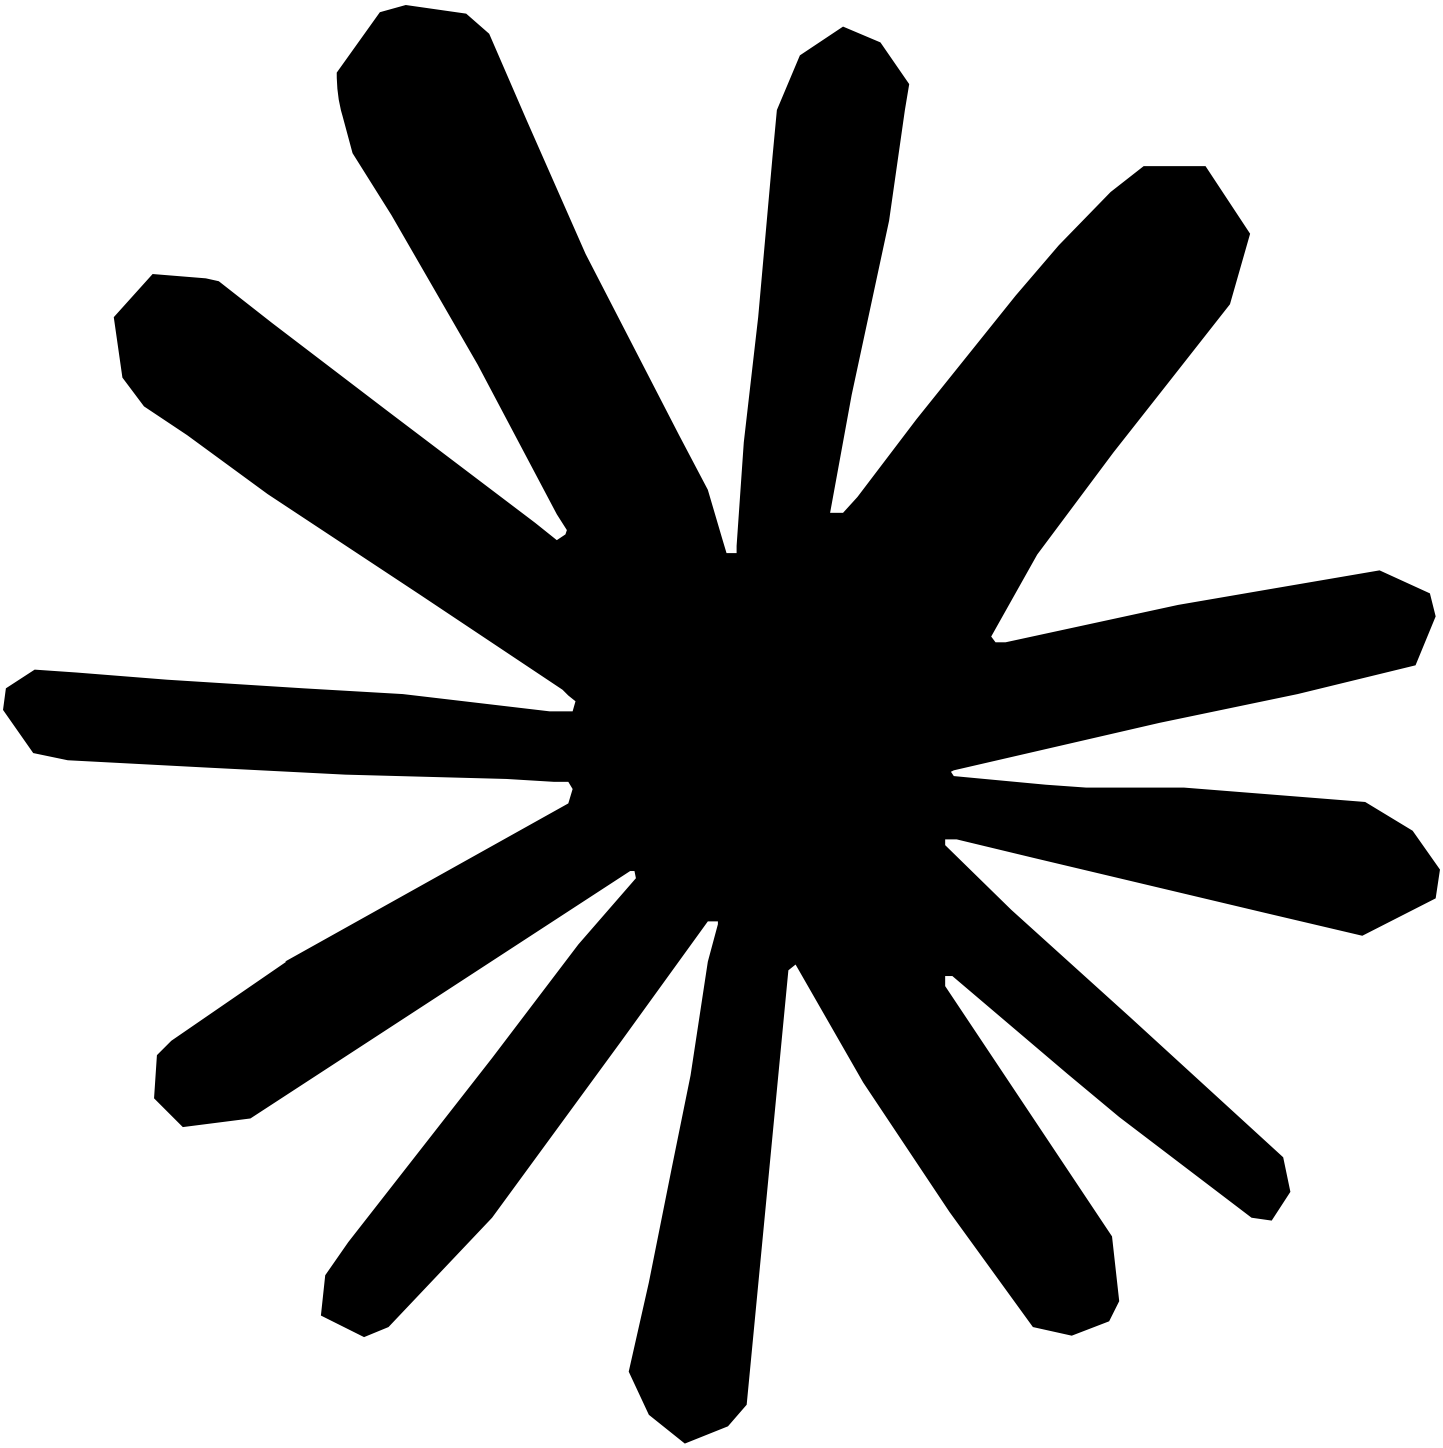

Claude can make mistakes.
Please double-check responses.

Choose files No file chosen

metadata: